



MySQL Database

MySQL is an open-source relational database management system (RDBMS) that uses Structured Query Language (SQL) to manage, manipulate, and retrieve data stored in databases. It is widely used for web applications and is part of the popular LAMP stack (Linux, Apache, MySQL, PHP/Perl/Python). MySQL is known for its ease of use, scalability, and performance.

MySQL is developed, distributed, and supported by Oracle Corporation. It is available in various editions, such as the community edition (free and open-source) and the enterprise edition (with additional features and support). MySQL provides a flexible and efficient way to store, organize, and manage data, making it suitable for a wide range of applications, from small personal projects to large-scale enterprise systems.

Introduction to SQL

SQL (Structured Query Language) is a standard language used to communicate with and manipulate relational databases such as MySQL, PostgreSQL, Oracle, and Microsoft SQL Server. With SQL, you can create, modify, and query data in databases, manage access control, and more.

Creating a Database

To create a new database, use the `CREATE DATABASE` statement followed by the database name.

Example:

```
CREATE DATABASE FrontlinesMediaDB;
```

This command creates a new database called 'FrontlinesMediaDB'. To work with this database, you need to select it using the `USE` statement:

```
USE FrontlinesMediaDB;
```

DDL (Data Definition Language)

DDL, or Data Definition Language, is a subset of SQL used to create, modify, and delete the structure of database objects such as tables, schemas, and views. DDL does not deal with data manipulation; it focuses on the structure of the database. In this section, we will explain DDL in detail with examples in MySQL.



1. CREATE TABLE

The CREATE TABLE statement is used to create a new table with columns, their data types, and constraints.

Example:

```
CREATE TABLE FLM_Articles (  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  title VARCHAR(255) NOT NULL,  
  author VARCHAR(255),  
  published_date DATE,  
  content TEXT  
);
```

This command creates a table called 'FLM_Articles' with five columns: 'id', 'title', 'author', 'published_date', and 'content'. The 'id' column is set to AUTO_INCREMENT and is defined as the PRIMARY KEY, ensuring that it has a unique value for each row.

2. ALTER TABLE

The ALTER TABLE statement is used to modify an existing table by adding, modifying, or deleting columns, as well as adding or dropping constraints.

Example - Add Column:

```
ALTER TABLE FLM_Articles  
ADD COLUMN category VARCHAR(255);
```

This command adds a new column called 'category' with a VARCHAR(255) data type to the existing 'FLM_Articles' table.

Example - Modify Column:

```
ALTER TABLE FLM_Articles  
MODIFY COLUMN author VARCHAR(100);
```

This command modifies the 'author' column in the 'FLM_Articles' table to have a VARCHAR(100) data type.

Example - Drop Column:

```
ALTER TABLE FLM_Articles  
DROP COLUMN content;
```

This command removes the 'content' column from the 'FLM_Articles' table.



3. RENAME:

The RENAME command is used to rename one or more tables. This operation is atomic, meaning it's performed in a single step and cannot be interrupted.

Example:

```
-- Create a table for FLM_Articles
CREATE TABLE FLM_Articles (
  id INT AUTO_INCREMENT PRIMARY KEY,
  title VARCHAR(255) NOT NULL,
  author VARCHAR(255) NOT NULL,
  content TEXT NOT NULL,
  published_date DATE NOT NULL
);

-- Rename the FLM_Articles table to FLM_BlogPosts
RENAME TABLE FLM_Articles TO FLM_BlogPosts;
```

In this example, the FLM_Articles table is created, and the RENAME command is used to rename it to FLM_BlogPosts.

These examples demonstrate how to use the TRUNCATE and RENAME commands in MySQL to delete all data from a table and rename a table, respectively.

4. TRUNCATE:

The TRUNCATE command is used to delete all data from a table without deleting the table structure itself. It is more efficient than the DELETE command when removing all data because it does not generate any undo logs or write any data to the binary log.

Example:

```
-- Create a table for FLM_Articles
CREATE TABLE FLM_Articles (
  id INT AUTO_INCREMENT PRIMARY KEY,
  title VARCHAR(255) NOT NULL,
  author VARCHAR(255) NOT NULL,
  content TEXT NOT NULL,
  published_date DATE NOT NULL
);

-- Insert data into the FLM_Articles table
INSERT INTO FLM_Articles (title, author, content, published_date)
```



```
VALUES ('MySQL Tutorial', 'John Doe', 'This is a MySQL tutorial...', '2023-05-02');
```

```
-- Truncate the FLM_Articles table, deleting all data  
TRUNCATE TABLE FLM_Articles;
```

In this example, the FLM_Articles table is created and populated with data. Then, the TRUNCATE command is used to delete all data from the table.

5. DROP TABLE

The DROP TABLE statement is used to delete an existing table and all its data.

Example:

```
DROP TABLE FLM_Articles;
```

This command deletes the 'FLM_Articles' table from the database.

6. CREATE INDEX

The CREATE INDEX statement is used to create an index on one or more columns of a table. Indexes can help speed up data retrieval operations.

Example:

```
CREATE INDEX idx_FLM_Articles_published_date ON FLM_Articles(published_date);
```

This command creates an index named 'idx_FLM_Articles_published_date' on the 'published_date' column of the 'FLM_Articles' table.

7. DROP INDEX

The DROP INDEX statement is used to delete an existing index.

Example:

```
DROP INDEX idx_FLM_Articles_published_date ON FLM_Articles;
```

This command deletes the 'idx_FLM_Articles_published_date' index from the 'FLM_Articles' table.

These examples provide a detailed understanding of DDL in MySQL, including creating tables, altering tables, dropping tables, creating indexes, and dropping indexes.



DML (Data Manipulation Language)

DML, or Data Manipulation Language, is a subset of SQL used to retrieve, insert, update, and delete data within database tables. It focuses on manipulating the data stored in a database rather than the structure of the database itself. In this section, we will explain DML in detail with examples in MySQL.

1. SELECT

The SELECT statement is used to retrieve data from one or more tables. You can specify the columns you want to retrieve, apply filters using the WHERE clause, and sort or group the results.

Example:

```
SELECT title, author, published_date FROM FLM_Articles WHERE category = 'Technology';
```

This command retrieves the 'title', 'author', and 'published_date' columns for all rows in the 'FLM_Articles' table where the 'category' is 'Technology'.

2. INSERT INTO

The INSERT INTO statement is used to insert new data into a table. You can specify the columns you want to insert data into and provide the data values.

Example:

```
INSERT INTO FLM_Articles (title, author, published_date, category)  
VALUES ('New FLM Tech Gadget', 'John Smith', '2023-05-05', 'Technology');
```

This command inserts a new row into the 'FLM_Articles' table with the values 'New FLM Tech Gadget' for 'title', 'John Smith' for 'author', '2023-05-05' for 'published_date', and 'Technology' for 'category'.

3. UPDATE

The UPDATE statement is used to modify existing data in a table. You can specify the new values for the columns and apply filters using the WHERE clause to determine which rows should be updated.

Example:

```
UPDATE FLM_Articles SET author = 'Jane Doe' WHERE id = 1;
```

This command updates the 'author' column to 'Jane Doe' for the row in the 'FLM_Articles' table where the 'id' is 1.



4. DELETE

The DELETE statement is used to remove data from a table. You can apply filters using the WHERE clause to determine which rows should be deleted.

Example:

```
DELETE FROM FLM_Articles WHERE id = 1;
```

This command deletes the row in the 'FLM_Articles' table where the 'id' is 1.

These examples provide a detailed understanding of DML in MySQL, including selecting, inserting, updating, and deleting data in tables.

DQL (Data Query Language)

DQL, or Data Query Language, is a subset of SQL used to retrieve data from databases, particularly from tables. It focuses on querying data rather than manipulating the data or the structure of the database itself. In this section, we will explain DQL in detail with examples in MySQL.

The primary DQL statement is the SELECT statement, which can be used with various clauses and functions to fetch data according to specific conditions, sorting, and grouping.

1. Basic SELECT

The simplest form of a SELECT statement is used to retrieve all columns from a table.

Example:

```
SELECT * FROM FLM_Articles;
```

This command retrieves all columns and rows from the 'FLM_Articles' table.

2. SELECT with specific columns

To retrieve only specific columns from a table, specify the column names after the SELECT keyword.

Example:

```
SELECT title, author, published_date FROM FLM_Articles;
```

This command retrieves the 'title', 'author', and 'published_date' columns for all rows in the 'FLM_Articles' table.

3. SELECT with WHERE clause

The WHERE clause is used to filter the results based on one or more conditions.



Example:

```
SELECT title, author, published_date FROM FLM_Articles WHERE category = 'Technology';
```

This command retrieves the 'title', 'author', and 'published_date' columns for all rows in the 'FLM_Articles' table where the 'category' is 'Technology'.

4. SELECT with ORDER BY clause

The ORDER BY clause is used to sort the results based on one or more columns, in either ascending (ASC) or descending (DESC) order.

Example:

```
SELECT title, author, published_date FROM FLM_Articles ORDER BY published_date DESC;
```

This command retrieves the 'title', 'author', and 'published_date' columns for all rows in the 'FLM_Articles' table and sorts the results by the 'published_date' column in descending order.

5. SELECT with GROUP BY and aggregate functions

The GROUP BY clause is used to group rows with the same values in specified columns. It is often used with aggregate functions like COUNT, SUM, AVG, MIN, or MAX to perform calculations on each group.

Example:

```
SELECT category, COUNT(*) as total_articles FROM FLM_Articles GROUP BY category;
```

This command retrieves the number of articles in each category in the 'FLM_Articles' table by grouping the rows by the 'category' column and counting the number of rows in each group.

These examples provide a detailed understanding of DQL in MySQL, including selecting data, applying filters with the WHERE clause, sorting results with the ORDER BY clause, and grouping data with the GROUP BY clause.

DCL (Data Control Language)

DCL, or Data Control Language, is a subset of SQL used to manage access control and permissions for database objects, such as tables, views, and procedures. It focuses on the security and authorization aspects of a database rather than data manipulation or the structure of the database itself. In this section, we will explain DCL in detail with examples in MySQL.



The primary DCL statements are GRANT and REVOKE, which are used to provide and remove access privileges to database objects.

1. GRANT

The GRANT statement is used to give specific privileges to a user or a group of users on a database object.

Example:

```
GRANT SELECT, INSERT, UPDATE ON FLM_Articles TO 'flm_user'@'localhost';
```

This command grants the 'flm_user' user the SELECT, INSERT, and UPDATE privileges on the 'FLM_Articles' table.

2. REVOKE

The REVOKE statement is used to remove specific privileges from a user or a group of users on a database object.

Example:

```
REVOKE INSERT, UPDATE ON FLM_Articles FROM 'flm_user'@'localhost';
```

This command revokes the INSERT and UPDATE privileges from the 'flm_user' user on the 'FLM_Articles' table.

It's important to note that the user management and privilege system in MySQL is quite extensive, and the examples provided here are a basic introduction to managing access control using DCL. In real-world scenarios, you may need to create and manage multiple users with different privileges, as well as manage access to multiple database objects.

These examples provide a detailed understanding of DCL in MySQL, including granting and revoking access privileges on database objects.

Example code for DDL, DML, DQL, and DCL:

```
-- DDL: Create a table for articles
CREATE TABLE FLM_Articles (
  id INT AUTO_INCREMENT PRIMARY KEY,
  title VARCHAR(255) NOT NULL,
  author VARCHAR(255) NOT NULL,
  content TEXT NOT NULL,
  published_date DATE NOT NULL
);
```

```
-- DML: Insert data into the articles table
```




```
INSERT INTO FLM_Articles (title, author, content, published_date)
VALUES ('MySQL Tutorial', 'John Doe', 'This is a MySQL tutorial...', '2023-05-02');
```

```
-- DQL: Select data from the articles table
SELECT * FROM FLM_Articles WHERE author = 'John Doe';
```

```
-- DCL: Grant permissions to a user on the articles table
GRANT SELECT, INSERT, UPDATE ON FLM_Articles TO 'flm_user'@'localhost';
```

In this example:

1. A DDL statement is used to create the FLM_Articles table, defining its structure with various columns and constraints.
2. A DML statement is used to insert a new row into the FLM_Articles table, providing values for the title, author, content, and published_date columns.
3. A DQL statement is used to query the FLM_Articles table, selecting all rows where the author is 'John Doe'.
4. A DCL statement is used to grant the SELECT, INSERT, and UPDATE permissions to a user called 'flm_user' on the FLM_Articles table.

This example demonstrates how to use DDL, DML, DQL, and DCL statements in MySQL to create, manipulate, query, and manage access control for a database table.

Constraints in MySQL

Constraints are rules applied to columns within a table to ensure data integrity and maintain the relationships between tables in a database. In this section, we will explain various constraints in MySQL, including Unique Key, Not Null, Primary Key, Composite Key, and Foreign Key.

1. UNIQUE Key

A UNIQUE constraint ensures that all values in a column are unique, meaning no two rows can have the same value in the specified column.

Example:

```
CREATE TABLE FLM_Authors (
  id INT AUTO_INCREMENT PRIMARY KEY,
  email VARCHAR(255) UNIQUE NOT NULL,
  name VARCHAR(255) NOT NULL
);
```



In this example, the 'email' column has a UNIQUE constraint, ensuring that each author has a unique email address.

2. NOT NULL

A NOT NULL constraint ensures that a column cannot have a NULL value, meaning a value must be provided for the specified column when a new row is inserted.

Example:

```
CREATE TABLE FLM_Articles (  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  title VARCHAR(255) NOT NULL,  
  author_id INT NOT NULL,  
  content TEXT NOT NULL  
);
```

In this example, the 'title', 'author_id', and 'content' columns have NOT NULL constraints, ensuring that these columns have values for every row in the table.

3. PRIMARY Key

A PRIMARY KEY constraint uniquely identifies each row in a table. It consists of one or more columns, and it enforces the uniqueness and NOT NULL constraints for the specified columns.

Example:

```
CREATE TABLE FLM_Articles (  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  title VARCHAR(255) NOT NULL,  
  author_id INT NOT NULL,  
  content TEXT NOT NULL  
);
```

In this example, the 'id' column is defined as the PRIMARY KEY, ensuring that it has a unique value for each row in the table.

4. Composite Key

A Composite Key is a key that consists of two or more columns to uniquely identify each row in a table. It is used when a single column is not sufficient to guarantee uniqueness.

Example:

```
CREATE TABLE FLM_Article_Tag_Relations (  
  article_id INT NOT NULL,  
  tag_id INT NOT NULL,  
  PRIMARY KEY (article_id, tag_id)
```



);

In this example, the combination of 'article_id' and 'tag_id' columns forms a Composite Key, ensuring that each row has a unique combination of both columns.

5. FOREIGN Key

A FOREIGN KEY constraint is used to create a relationship between two tables. It enforces referential integrity by ensuring that the values in a column of one table match the values in the primary key column of another table.

Example:

```
CREATE TABLE FLM_Articles (  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  title VARCHAR(255) NOT NULL,  
  author_id INT NOT NULL,  
  content TEXT NOT NULL,  
  FOREIGN KEY (author_id) REFERENCES FLM_Authors(id)  
);
```

In this example, the 'author_id' column in the 'FLM_Articles' table is defined as a FOREIGN KEY that references the 'id' column in the 'FLM_Authors' table. This ensures that every 'author_id' value in the 'FLM_Articles' table corresponds to a valid 'id' value in the 'FLM_Authors' table.

These examples provide a detailed understanding of constraints in MySQL, including Unique Key, Not Null, Primary Key, Composite Key, and Foreign Key.

Example code for Constraints:

In this example, we will create a simple EdTech database schema to manage courses, lessons, and students using MySQL. We will incorporate various constraints, including Unique Key, Not Null, Primary Key, Composite Key, and Foreign Key.

-- Create a table for students

```
CREATE TABLE FLM_Students (  
  student_id INT AUTO_INCREMENT PRIMARY KEY,  
  email VARCHAR(255) UNIQUE NOT NULL,  
  first_name VARCHAR(255) NOT NULL,  
  last_name VARCHAR(255) NOT NULL  
);
```

-- Create a table for courses

```
CREATE TABLE FLM_Courses (  

```



```
course_id INT AUTO_INCREMENT PRIMARY KEY,  
course_name VARCHAR(255) UNIQUE NOT NULL,  
course_description TEXT NOT NULL  
);  
  
-- Create a table for lessons  
CREATE TABLE FLM_Lessons (  
    lesson_id INT AUTO_INCREMENT PRIMARY KEY,  
    course_id INT NOT NULL,  
    lesson_name VARCHAR(255) NOT NULL,  
    lesson_content TEXT NOT NULL,  
    FOREIGN KEY (course_id) REFERENCES FLM_Courses(course_id)  
);  
  
-- Create a table for student_course_relations  
CREATE TABLE FLM_Student_Course_Relations (  
    student_id INT NOT NULL,  
    course_id INT NOT NULL,  
    enrollment_date DATE NOT NULL,  
    PRIMARY KEY (student_id, course_id),  
    FOREIGN KEY (student_id) REFERENCES FLM_Students(student_id),  
    FOREIGN KEY (course_id) REFERENCES FLM_Courses(course_id)  
);
```

In this example:

1. The FLM_Students table stores student information, with a PRIMARY KEY constraint on the student_id column and a UNIQUE constraint on the email column.
2. The FLM_Courses table stores course information, with a PRIMARY KEY constraint on the course_id column and a UNIQUE constraint on the course_name column.
3. The FLM_Lessons table stores lesson information, with a PRIMARY KEY constraint on the lesson_id column and a FOREIGN KEY constraint on the course_id column, referencing the course_id column in the FLM_Courses table.
4. The FLM_Student_Course_Relations table stores the relationships between students and courses, with a Composite PRIMARY KEY consisting of the student_id and course_id columns, and FOREIGN KEY constraints on both the student_id and course_id columns, referencing the student_id column in the FLM_Students table and the course_id column in the FLM_Courses table, respectively.



This EdTech example demonstrates how to create a MySQL database schema using various constraints.

Aggregate functions

Aggregate functions in MySQL are used to perform calculations on a set of values and return a single value. They are often used with the GROUP BY clause to group rows that share a common property and perform calculations on each group. In this section, we will explain the MIN, MAX, AVG, SUM, and COUNT aggregate functions in MySQL.

1. MIN:

The MIN function returns the minimum value of a specified column.

Example:

```
CREATE TABLE FLM_Products (  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  name VARCHAR(255) NOT NULL,  
  price DECIMAL(10, 2) NOT NULL  
);
```

```
-- Find the minimum price among all products  
SELECT MIN(price) AS min_price FROM FLM_Products;
```

2. MAX:

The MAX function returns the maximum value of a specified column.

Example:

```
-- Find the maximum price among all products  
SELECT MAX(price) AS max_price FROM FLM_Products;
```

3. AVG:

The AVG function returns the average value of a specified column.

Example:

```
-- Find the average price of all products  
SELECT AVG(price) AS avg_price FROM FLM_Products;
```

4. SUM:

The SUM function returns the sum of all values in a specified column.

**Example:**

-- Find the total sum of prices for all products

```
SELECT SUM(price) AS total_price FROM FLM_Products;
```

5. COUNT:

The COUNT function returns the number of rows that match a specified condition. When used with an asterisk (*), it returns the total number of rows in the table.

Example:

-- Count the total number of products

```
SELECT COUNT(*) AS product_count FROM FLM_Products;
```

-- Count the number of products with a price greater than 50

```
SELECT COUNT(*) AS expensive_product_count FROM FLM_Products WHERE price > 50;
```

- The above examples provide a detailed understanding of the MIN, MAX, AVG, SUM, and COUNT aggregate functions in MySQL.

Example code for all aggregate functions:

In this example, we will create a table called FLM_Sales to store sales data and demonstrate the use of all the aggregate functions: MIN, MAX, AVG, SUM, and COUNT.

-- Create the FLM_Sales table

```
CREATE TABLE FLM_Sales (  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  sales_person VARCHAR(255) NOT NULL,  
  sale_amount DECIMAL(10, 2) NOT NULL,  
  sale_date DATE NOT NULL  
);
```

-- Insert sample data into the FLM_Sales table

```
INSERT INTO FLM_Sales (sales_person, sale_amount, sale_date)  
VALUES ('John Doe', 500.00, '2023-01-01'),  
      ('John Doe', 750.00, '2023-01-02'),  
      ('Jane Smith', 1200.00, '2023-01-03'),  
      ('Jane Smith', 650.00, '2023-01-04'),  
      ('John Doe', 900.00, '2023-01-05');
```



```
-- Use aggregate functions to analyze the data
SELECT
  MIN(sale_amount) AS min_sale_amount,
  MAX(sale_amount) AS max_sale_amount,
  AVG(sale_amount) AS avg_sale_amount,
  SUM(sale_amount) AS total_sale_amount,
  COUNT(*) AS total_sales
FROM FLM_Sales;
```

In this example, we first create the FLM_Sales table and insert sample data. Then, we use a single SELECT statement to perform calculations using all the aggregate functions:

1. The MIN function returns the minimum sale amount.
2. The MAX function returns the maximum sale amount.
3. The AVG function returns the average sale amount.
4. The SUM function returns the total sale amount.
5. The COUNT function returns the total number of sales.

The above example demonstrates how to use multiple aggregate functions in a single query in MySQL.

Joins in MySQL

SQL joins in MySQL are used to combine rows from two or more tables based on a related column between them. There are several types of SQL joins: INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN. MySQL does not support FULL OUTER JOIN directly, but it can be emulated. We will use variable names related to Frontlines Media (FLM).

1. INNER JOIN:

The INNER JOIN keyword selects records that have matching values in both tables.

Example:

```
CREATE TABLE FLM_Authors (
  id INT AUTO_INCREMENT PRIMARY KEY,
  name VARCHAR(255) NOT NULL
);
```



```
CREATE TABLE FLM_Articles (  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  author_id INT NOT NULL,  
  title VARCHAR(255) NOT NULL,  
  FOREIGN KEY (author_id) REFERENCES FLM_Authors(id)  
);
```

```
-- Select all articles and their authors using INNER JOIN  
SELECT FLM_Authors.name, FLM_Articles.title  
FROM FLM_Authors  
INNER JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id;
```

2. LEFT JOIN (or LEFT OUTER JOIN):

The LEFT JOIN keyword returns all records from the left table (FLM_Authors), and the matched records from the right table (FLM_Articles). If no match is found, NULL values are returned for the right table's columns.

Example:

```
-- Select all authors and their articles (if any) using LEFT JOIN  
SELECT FLM_Authors.name, FLM_Articles.title  
FROM FLM_Authors  
LEFT JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id;
```

3. RIGHT JOIN (or RIGHT OUTER JOIN):

The RIGHT JOIN keyword returns all records from the right table (FLM_Articles), and the matched records from the left table (FLM_Authors). If no match is found, NULL values are returned for the left table's columns.

Example:

```
-- Select all articles and their authors (if any) using RIGHT JOIN  
SELECT FLM_Authors.name, FLM_Articles.title  
FROM FLM_Authors  
RIGHT JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id;
```

4. FULL OUTER JOIN:

MySQL does not support FULL OUTER JOIN directly. However, it can be emulated using a combination of LEFT JOIN and RIGHT JOIN with a UNION clause. The FULL OUTER JOIN keyword returns all records when there is a match in either the left or right table.

**Example:**

```
-- Emulate FULL OUTER JOIN in MySQL
SELECT FLM_Authors.name, FLM_Articles.title
FROM FLM_Authors
LEFT JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id
UNION
SELECT FLM_Authors.name, FLM_Articles.title
FROM FLM_Authors
RIGHT JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id;
```

These examples provide a detailed understanding of the different types of SQL joins in MySQL.

Example code for Joins:

In this example, we will create two tables, FLM_Authors and FLM_Articles, and demonstrate the use of all SQL joins in MySQL: INNER JOIN, LEFT JOIN, RIGHT JOIN, and emulated FULL OUTER JOIN.

```
-- Create the FLM_Authors table
CREATE TABLE FLM_Authors (
  id INT AUTO_INCREMENT PRIMARY KEY,
  name VARCHAR(255) NOT NULL
);
```

```
-- Create the FLM_Articles table
CREATE TABLE FLM_Articles (
  id INT AUTO_INCREMENT PRIMARY KEY,
  author_id INT NOT NULL,
  title VARCHAR(255) NOT NULL,
  FOREIGN KEY (author_id) REFERENCES FLM_Authors(id)
);
```

```
-- Insert sample data into FLM_Authors and FLM_Articles tables
INSERT INTO FLM_Authors (name) VALUES ('John Doe'), ('Jane Smith'), ('Mark Johnson');
INSERT INTO FLM_Articles (author_id, title) VALUES (1, 'MySQL Tutorial'), (1, 'Python Tutorial'), (3, 'JavaScript Tutorial');
```

```
-- INNER JOIN example
```



```
SELECT FLM_Authors.name, FLM_Articles.title
FROM FLM_Authors
INNER JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id;
```

-- LEFT JOIN example

```
SELECT FLM_Authors.name, FLM_Articles.title
FROM FLM_Authors
LEFT JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id;
```

-- RIGHT JOIN example

```
SELECT FLM_Authors.name, FLM_Articles.title
FROM FLM_Authors
RIGHT JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id;
```

-- FULL OUTER JOIN emulation in MySQL

```
SELECT FLM_Authors.name, FLM_Articles.title
FROM FLM_Authors
LEFT JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id
UNION
SELECT FLM_Authors.name, FLM_Articles.title
FROM FLM_Authors
RIGHT JOIN FLM_Articles ON FLM_Authors.id = FLM_Articles.author_id;
```

This example demonstrates how to use all types of SQL joins in MySQL with the FLM_Authors and FLM_Articles tables. The INNER JOIN, LEFT JOIN, and RIGHT JOIN examples return the author names with their corresponding article titles. The emulated FULL OUTER JOIN returns all records from both tables, filling in NULL values where no match is found.



JDBC

Attribute	Description
What is JDBC?	JDBC (Java Database Connectivity) is an API (Application Programming Interface) for the Java programming language that defines how a client may access a database. It provides methods to query, insert, update, and delete data in a relational database, and allows Java applications to interact with various databases using a standard interface, regardless of the underlying database management system (DBMS).
Applications	1. Web applications: JDBC is commonly used in Java-based web applications to interact with databases for storing and retrieving data, such as user information, product details, and orders.
	2. Enterprise applications: Java Enterprise Edition (Java EE) applications use JDBC to integrate with databases for handling complex business transactions and data processing.
	3. Desktop applications: Java desktop applications, such as management systems and data analysis tools, use JDBC to connect to databases for data storage and retrieval.
	4. Data migration: JDBC can be used to migrate data between different databases or different versions of the same database.
	5. Reporting tools: Java-based reporting tools, like JasperReports or BIRT, use JDBC to fetch data from databases for generating reports.
Why is it used?	1. Platform independence: JDBC provides a consistent and platform-independent way to interact with databases, allowing developers to write database applications that can run on various platforms without modification.



	<p>2. Database compatibility: JDBC allows Java applications to connect to a wide range of databases using a standard interface, which makes it easier to switch between different DBMSs without changing the application code.</p>
	<p>3. Ease of use: JDBC provides a simple and intuitive API that makes it easy for developers to perform common database operations like creating, reading, updating, and deleting records.</p>
	<p>4. Flexibility: JDBC supports various types of databases (relational, object-relational, etc.) and different data sources (local, remote, cloud-based), providing developers with flexibility in choosing the most suitable database solution for their applications.</p>
	<p>5. Integration with Java ecosystem: As part of the Java platform, JDBC integrates seamlessly with other Java technologies, such as Java servlets, JavaServer Pages (JSP), JavaServer Faces (JSF), and Java Persistence API (JPA), making it easier to build comprehensive and scalable applications.</p>
Scope	<p>The scope of JDBC extends to any Java-based application or system that requires interaction with databases. This encompasses a wide range of applications, from small-scale desktop applications to large-scale enterprise systems. JDBC is compatible with various database management systems (DBMS), such as MySQL, PostgreSQL, Oracle, Microsoft SQL Server, and many others.</p>

SQLExceptions

SQLExceptions are exceptions that occur when an error is encountered during interaction with a database using JDBC. SQLExceptions provide information about database access errors or other errors related to JDBC. They usually occur due to incorrect SQL syntax, issues with the database connection, or other problems during database operations.



Examples of common SQLExceptions include:

1. Incorrect SQL syntax
2. Invalid database URL, username, or password
3. Insufficient privileges to perform a specific operation
4. Violation of database constraints like foreign keys or unique keys

To handle SQLExceptions, you should use a try-catch block when performing JDBC operations. This allows you to catch any SQLExceptions that occur and provide an appropriate error message or take any necessary action to handle the error.

Here's an example demonstrating how to handle SQLExceptions using a Java application with JDBC:

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.ResultSet;
import java.sql.SQLException;
import java.sql.Statement;

public class FLM_JdbcExample {
    public static void main(String[] args) {
        Connection connection = null;
        Statement statement = null;
        ResultSet resultSet = null;

        try {
            Class.forName("com.mysql.cj.jdbc.Driver");
            String url = "jdbc:mysql://localhost:3306/FLM_DB";
            String user = "root";
            String password = "your_password";
            connection = DriverManager.getConnection(url, user, password);

            statement = connection.createStatement();
            String sql = "SELECT * FROM FLM_Authors";
            resultSet = statement.executeQuery(sql);

            while (resultSet.next()) {
                int id = resultSet.getInt("id");
                String name = resultSet.getString("name");
            }
        }
    }
}
```



```
        System.out.println("ID: " + id + ", Name: " + name);
    }

    } catch (ClassNotFoundException e) {
        System.out.println("Error loading JDBC driver: " + e.getMessage());
        e.printStackTrace();
    } catch (SQLException e) {
        System.out.println("Error executing SQL query: " + e.getMessage());
        e.printStackTrace();
    } finally {
        try {
            if (resultSet != null) resultSet.close();
            if (statement != null) statement.close();
            if (connection != null) connection.close();
        } catch (SQLException e) {
            System.out.println("Error closing resources: " + e.getMessage());
            e.printStackTrace();
        }
    }
}
}
```

In this example, we use a try-catch block to handle `ClassNotFoundException` and `SQLException`s. If an `SQLException` occurs, it will be caught and an appropriate error message will be displayed. Additionally, we use a finally block to close the resources after the operation is complete, ensuring that any `SQLException`s that occur during closing are also caught and handled.

DriverManager

In MySQL, `DriverManager` is a class in the JDBC API that manages the registered JDBC drivers. It provides a method to establish a connection to a database using a JDBC driver. The `DriverManager` class is responsible for finding the appropriate driver for the specific database URL provided and loading it.

To use the `DriverManager` in MySQL with JDBC, you must first load the driver class. Once the driver is loaded, you can obtain a `Connection` object to connect to the database using the `DriverManager.getConnection()` method.



Here's an example demonstrating how to use DriverManager to connect to a MySQL database using JDBC:

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.SQLException;

public class FLM_JdbcExample {
    public static void main(String[] args) {
        Connection connection = null;
        try {
            Class.forName("com.mysql.cj.jdbc.Driver");
            String url = "jdbc:mysql://localhost:3306/FLM_DB";
            String user = "root";
            String password = "your_password";
            connection = DriverManager.getConnection(url, user, password);
            System.out.println("Connected to the FLM_DB database successfully!");

            // Perform database operations using the connection object...

        } catch (ClassNotFoundException e) {
            System.out.println("Error loading JDBC driver: " + e.getMessage());
            e.printStackTrace();
        } catch (SQLException e) {
            System.out.println("Error connecting to database: " + e.getMessage());
            e.printStackTrace();
        } finally {
            try {
                if (connection != null) connection.close();
            } catch (SQLException e) {
                System.out.println("Error closing connection: " + e.getMessage());
                e.printStackTrace();
            }
        }
    }
}
```

In this example, we first load the MySQL JDBC driver class using the `Class.forName()` method. Then we provide the database URL, username, and password to



`DriverManager.getConnection()` to create a `Connection` object. If the connection is successful, we print a message indicating that we have connected to the database.

If there is an exception thrown by the `Class.forName()` or `DriverManager.getConnection()` methods, the corresponding catch block will execute, and the error message will be printed. We also use a finally block to close the `Connection` object to ensure that resources are freed up and to handle any `SQLExceptions` that may occur during the closing process.

Overall, `DriverManager` provides a simple and convenient way to manage and establish a connection to a MySQL database using JDBC.

Connection

In MySQL, a `Connection` object represents a physical connection to a database. It's created using the `DriverManager.getConnection()` method and is used to interact with the database.

The `Connection` object provides methods to execute SQL queries, commit or rollback transactions, and manage session-level parameters like auto-commit and transaction isolation level.

Here's an example demonstrating how to create a `Connection` object and execute a simple SQL query using JDBC:

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.ResultSet;
import java.sql.SQLException;
import java.sql.Statement;

public class FLM_JdbcExample {
    public static void main(String[] args) {
        Connection connection = null;
        Statement statement = null;
        ResultSet resultSet = null;

        try {
            Class.forName("com.mysql.cj.jdbc.Driver");
```




```
String url = "jdbc:mysql://localhost:3306/FLM_DB";
String user = "root";
String password = "your_password";
connection = DriverManager.getConnection(url, user, password);
System.out.println("Connected to the FLM_DB database successfully!");

statement = connection.createStatement();
String sql = "SELECT * FROM FLM_Authors";
resultSet = statement.executeQuery(sql);

while (resultSet.next()) {
    int id = resultSet.getInt("id");
    String name = resultSet.getString("name");
    System.out.println("ID: " + id + ", Name: " + name);
}

} catch (ClassNotFoundException e) {
    System.out.println("Error loading JDBC driver: " + e.getMessage());
    e.printStackTrace();
} catch (SQLException e) {
    System.out.println("Error connecting to database: " + e.getMessage());
    e.printStackTrace();
} finally {
    try {
        if (resultSet != null) resultSet.close();
        if (statement != null) statement.close();
        if (connection != null) connection.close();
    } catch (SQLException e) {
        System.out.println("Error closing resources: " + e.getMessage());
        e.printStackTrace();
    }
}
}
```

In this example, we first create a Connection object using the DriverManager.getConnection() method. We then create a Statement object using the Connection.createStatement() method and execute a SQL query to select all records from the FLM_Authors table. We then iterate through the ResultSet object and print the results to the console.



After the operation is complete, we use a finally block to close the ResultSet, Statement, and Connection objects to ensure that resources are freed up and to handle any SQLExceptions that may occur during the closing process.

Overall, the Connection object is a key component in JDBC as it provides a way to interact with the database and manage transactions.

Statement

In MySQL, a Statement object represents a SQL statement that is executed against a database. It's created using the Connection.createStatement() method and is used to execute SQL queries and updates.

There are three types of Statement objects in MySQL:

1. Statement: Used for executing simple SQL statements without any parameters.
2. PreparedStatement: Used for executing SQL statements that have parameters.
3. CallableStatement: Used for executing stored procedures.

Statement objects provide methods to execute SQL queries, retrieve data from the database, and update the database.

Here's an example demonstrating how to create a Statement object and execute a simple SQL query using JDBC:

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.ResultSet;
import java.sql.SQLException;
import java.sql.Statement;

public class FLM_JdbcExample {
    public static void main(String[] args) {
        Connection connection = null;
        Statement statement = null;
        ResultSet resultSet = null;

        try {
            Class.forName("com.mysql.cj.jdbc.Driver");
```



```
String url = "jdbc:mysql://localhost:3306/FLM_DB";
String user = "root";
String password = "your_password";
connection = DriverManager.getConnection(url, user, password);
System.out.println("Connected to the FLM_DB database successfully!");

statement = connection.createStatement();
String sql = "SELECT * FROM FLM_Authors";
resultSet = statement.executeQuery(sql);

while (resultSet.next()) {
    int id = resultSet.getInt("id");
    String name = resultSet.getString("name");
    System.out.println("ID: " + id + ", Name: " + name);
}

} catch (ClassNotFoundException e) {
    System.out.println("Error loading JDBC driver: " + e.getMessage());
    e.printStackTrace();
} catch (SQLException e) {
    System.out.println("Error executing SQL statement: " + e.getMessage());
    e.printStackTrace();
} finally {
    try {
        if (resultSet != null) resultSet.close();
        if (statement != null) statement.close();
        if (connection != null) connection.close();
    } catch (SQLException e) {
        System.out.println("Error closing resources: " + e.getMessage());
        e.printStackTrace();
    }
}
}
```

In this example, we first create a Connection object using the DriverManager.getConnection() method. We then create a Statement object using the Connection.createStatement() method and execute a SQL query to select all records from the FLM_Authors table. We then iterate through the ResultSet object and print the results to the console.



After the operation is complete, we use a finally block to close the ResultSet, Statement, and Connection objects to ensure that resources are freed up and to handle any SQLExceptions that may occur during the closing process.

Overall, the Statement object is a key component in JDBC as it provides a way to execute SQL statements against a database and retrieve or update data.

PreparedStatement

In MySQL, a PreparedStatement object is a subclass of Statement that represents a precompiled SQL statement with parameters. It's used when the same SQL statement is executed multiple times with different parameter values.

PreparedStatement objects provide a way to reduce SQL injection attacks and improve performance by allowing the database to reuse the execution plan for the same statement.

Here's an example demonstrating how to create a PreparedStatement object and execute a parameterized SQL query using JDBC:

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.PreparedStatement;
import java.sql.ResultSet;
import java.sql.SQLException;

public class FLM_JdbcExample {
    public static void main(String[] args) {
        Connection connection = null;
        PreparedStatement preparedStatement = null;
        ResultSet resultSet = null;

        try {
            Class.forName("com.mysql.cj.jdbc.Driver");
            String url = "jdbc:mysql://localhost:3306/FLM_DB";
            String user = "root";
            String password = "your_password";
```



```
connection = DriverManager.getConnection(url, user, password);
System.out.println("Connected to the FLM_DB database successfully!");

String sql = "SELECT * FROM FLM_Authors WHERE id=?";
PreparedStatement = connection.prepareStatement(sql);
PreparedStatement.setInt(1, 1); // Set the parameter value

resultSet = PreparedStatement.executeQuery();

while (resultSet.next()) {
    int id = resultSet.getInt("id");
    String name = resultSet.getString("name");
    System.out.println("ID: " + id + ", Name: " + name);
}

} catch (ClassNotFoundException e) {
    System.out.println("Error loading JDBC driver: " + e.getMessage());
    e.printStackTrace();
} catch (SQLException e) {
    System.out.println("Error executing SQL statement: " + e.getMessage());
    e.printStackTrace();
} finally {
    try {
        if (resultSet != null) resultSet.close();
        if (PreparedStatement != null) PreparedStatement.close();
        if (connection != null) connection.close();
    } catch (SQLException e) {
        System.out.println("Error closing resources: " + e.getMessage());
        e.printStackTrace();
    }
}
}
```

In this example, we first create a Connection object using the DriverManager.getConnection() method. We then create a PreparedStatement object using the Connection.prepareStatement() method and execute a parameterized SQL query to select records from the FLM_Authors table where the id column equals 1. We then iterate through the ResultSet object and print the results to the console.



After the operation is complete, we use a finally block to close the ResultSet, PreparedStatement, and Connection objects to ensure that resources are freed up and to handle any SQLExceptions that may occur during the closing process.

Overall, the PreparedStatement object is a key component in JDBC as it provides a way to execute parameterized SQL queries and improve performance by reusing the execution plan for the same statement.

ResultSet

In MySQL, a ResultSet object represents a set of results returned by a SQL query. It's created by executing a query using a Statement or PreparedStatement object and provides methods to iterate through the rows of the result set and retrieve data from each column.

ResultSet objects are used to retrieve data from the database and are usually obtained by executing a SQL query. They provide methods to move the cursor to the next row, get the value of a column in the current row, and navigate to a specific row or column in the result set.

Here's an example demonstrating how to create a ResultSet object and iterate through the rows of the result set using JDBC:

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.ResultSet;
import java.sql.SQLException;
import java.sql.Statement;

public class FLM_JdbcExample {
    public static void main(String[] args) {
        Connection connection = null;
        Statement statement = null;
        ResultSet resultSet = null;

        try {
            Class.forName("com.mysql.cj.jdbc.Driver");
            String url = "jdbc:mysql://localhost:3306/FLM_DB";
```



```
String user = "root";
String password = "your_password";
connection = DriverManager.getConnection(url, user, password);
System.out.println("Connected to the FLM_DB database successfully!");

statement = connection.createStatement();
String sql = "SELECT * FROM FLM_Authors";
resultSet = statement.executeQuery(sql);

while (resultSet.next()) {
    int id = resultSet.getInt("id");
    String name = resultSet.getString("name");
    System.out.println("ID: " + id + ", Name: " + name);
}

} catch (ClassNotFoundException e) {
    System.out.println("Error loading JDBC driver: " + e.getMessage());
    e.printStackTrace();
} catch (SQLException e) {
    System.out.println("Error executing SQL statement: " + e.getMessage());
    e.printStackTrace();
} finally {
    try {
        if (resultSet != null) resultSet.close();
        if (statement != null) statement.close();
        if (connection != null) connection.close();
    } catch (SQLException e) {
        System.out.println("Error closing resources: " + e.getMessage());
        e.printStackTrace();
    }
}
}
```

In this example, we first create a Connection object using the DriverManager.getConnection() method. We then create a Statement object using the Connection.createStatement() method and execute a SQL query to select all records from the FLM_Authors table. We then iterate through the ResultSet object using the next() method, which moves the cursor to the next row and returns true if there is a next row, or false if there are no more rows.



Inside the loop, we use the `getInt()` and `getString()` methods to retrieve the values of the `id` and `name` columns in the current row, respectively.

After the operation is complete, we use a `finally` block to close the `ResultSet`, `Statement`, and `Connection` objects to ensure that resources are freed up and to handle any `SQLExceptions` that may occur during the closing process.

Overall, the `ResultSet` object is a key component in JDBC as it provides a way to retrieve and manipulate data from the database.

CRUD application using JDBC

A CRUD application using JDBC in MySQL is an application that performs Create, Read, Update, and Delete operations on a database using JDBC. The application connects to a MySQL database using JDBC, executes SQL statements to perform CRUD operations on the database, and retrieves and displays data to the user.

Here's an example demonstrating how to create a simple CRUD application using JDBC in MySQL:

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.PreparedStatement;
import java.sql.ResultSet;
import java.sql.SQLException;
import java.sql.Statement;

public class FLM_JdbcExample {
    public static void main(String[] args) {
        Connection connection = null;
        Statement statement = null;
        PreparedStatement preparedStatement = null;
        ResultSet resultSet = null;

        try {
            Class.forName("com.mysql.cj.jdbc.Driver");
            String url = "jdbc:mysql://localhost:3306/FLM_DB";
```




```
String user = "root";
String password = "your_password";
connection = DriverManager.getConnection(url, user, password);
System.out.println("Connected to the FLM_DB database successfully!");

// CREATE
String insertSql = "INSERT INTO FLM_Books (title, author, price) VALUES (?, ?, ?)";
preparedStatement = connection.prepareStatement(insertSql);
preparedStatement.setString(1, "The Lord of the Rings");
preparedStatement.setString(2, "J.R.R. Tolkien");
preparedStatement.setDouble(3, 20.99);
preparedStatement.executeUpdate();
System.out.println("New book added to the database!");

// READ
statement = connection.createStatement();
String selectSql = "SELECT * FROM FLM_Books";
resultSet = statement.executeQuery(selectSql);

while (resultSet.next()) {
    int id = resultSet.getInt("id");
    String title = resultSet.getString("title");
    String author = resultSet.getString("author");
    double price = resultSet.getDouble("price");
    System.out.println("ID: " + id + ", Title: " + title + ", Author: " + author + ", Price: " +
price);
}

// UPDATE
String updateSql = "UPDATE FLM_Books SET price=? WHERE title=?";
preparedStatement = connection.prepareStatement(updateSql);
preparedStatement.setDouble(1, 24.99);
preparedStatement.setString(2, "The Lord of the Rings");
preparedStatement.executeUpdate();
System.out.println("Book price updated in the database!");

// DELETE
String deleteSql = "DELETE FROM FLM_Books WHERE title=?";
preparedStatement = connection.prepareStatement(deleteSql);
preparedStatement.setString(1, "The Lord of the Rings");
```



```
preparedStatement.executeUpdate();
System.out.println("Book deleted from the database!");

} catch (ClassNotFoundException e) {
    System.out.println("Error loading JDBC driver: " + e.getMessage());
    e.printStackTrace();
} catch (SQLException e) {
    System.out.println("Error executing SQL statement: " + e.getMessage());
    e.printStackTrace();
} finally {
    try {
        if (resultSet != null) resultSet.close();
        if (statement != null) statement.close();
        if (preparedStatement != null) preparedStatement.close();
        if (connection != null) connection.close();
    } catch (SQLException e) {
        System.out.println("Error closing resources: " + e.getMessage());
        e.printStackTrace();
    }
}
}
```

In this example, we first create a Connection object using the DriverManager.getConnection() method. We then demonstrate all four CRUD operations:

1. Create: We create a new book record by using a PreparedStatement object to execute an INSERT statement.
2. Read: We retrieve all book records from the FLM_Books table using a Statement object to execute a SELECT statement and iterate through the ResultSet object to retrieve the values of the columns.
3. Update: We update the price of a book record by using a PreparedStatement object to execute an UPDATE statement.
4. Delete: We delete a book record by using a PreparedStatement object to execute a DELETE statement.

After each operation, we display a message to the user indicating whether the operation was successful or not.



In this example, we use both Statement and PreparedStatement objects to execute SQL statements. The difference between the two is that PreparedStatement objects are precompiled, which makes them more efficient when executing the same statement multiple times with different parameter values.

Overall, a CRUD application using JDBC in MySQL is a useful way to interact with a database and perform basic database operations. It's an essential part of any database-driven application and provides a way to manage the data in the database.